

pamica: GPU-accelerated Adaptive Mixture Independent Component Analysis in Python with Fortran parity

Seyed Yahya Shirazi¹, Arnaud Delorme^{1,2}, and Scott Makeig¹

¹ Swartz Center for Computational Neuroscience, Institute for Neural Computation, University of California San Diego, USA ² Centre de Recherche Cerveau et Cognition (CerCo), CNRS, University of Toulouse, France ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Independent Component Analysis (ICA) is a standard method for separating electroencephalography (EEG) and electromyography (EMG) recordings into maximally independent sources that isolate brain, muscle, and artifact activity for downstream analysis. Adaptive Mixture ICA (AMICA) ([Palmer et al., 2012](#)) generalizes single-model ICA to a mixture of such models with adaptive source densities, and produces the most dipolar (and thus most physiologically interpretable) component decompositions of EEG among the widely used algorithms benchmarked by Delorme et al. ([2012](#)). Its reference implementation is a Fortran program parallelized with the Message Passing Interface (MPI) and distributed as a compiled binary driven from MATLAB/EEGLAB, which is difficult to install, runs only on the central processing unit (CPU), and is not usable from a Python scientific workflow.

pamica is a Python implementation of AMICA that reproduces the reference Fortran results within numerical tolerance while running on the CPU, NVIDIA graphics processing units (GPUs, via CUDA), and Apple GPUs (Apple's MLX array framework ([Hannun et al., 2023](#))). It is a complete reimplement built on PyTorch ([Paszke et al., 2019](#)), NumPy ([Harris et al., 2020](#)), and SciPy ([Virtanen et al., 2020](#)), not a wrapper around the Fortran binary, and exposes a scikit-learn-style estimator under a BSD-3-Clause license. In double precision it reproduces the reference score-function algebra to floating-point round-off; it also runs in single precision (float32), unavailable in the CPU-only binary and required for Apple GPUs, which have no float64 and host the fastest backend (MLX). pamica writes output in the binary format that EEGLAB's AMICA loader reads: a single-model output file is byte-identical in layout to a native AMICA file, and multi-model output round-trips through the same loader. Correctness is defined as parity with the Fortran reference for the single-model case and, because multi-model AMICA is not partition-identifiable, as a similar distribution of solutions for the multi-model case; both are validated on real EEG against the reference binary. The software is at <https://github.com/sccn/pAMICA> (archived at doi:10.5281/zenodo.21312148).

Statement of need

AMICA's decompositions are well suited to equivalent-dipole source localization and automated component classification ([Pion-Tonachini et al., 2019](#)). Yet its reference implementation is MATLAB-only Fortran, an increasing obstacle as neuroimaging analysis moves toward Python, for example MNE-Python ([Gramfort et al., 2013](#)): an AMICA that runs natively in Python and on a GPU, and that is validated to reproduce the Fortran reference numerically, is needed for modern pipelines.

General-purpose Python ICA implementations do not fill this gap. scikit-learn and MNE-Python provide FastICA (Hyvärinen & Oja, 2000) and Infomax (Bell & Sejnowski, 1995; Lee et al., 1999), while Picard (Ablin et al., 2018) offers faster-converging maximum-likelihood ICA; none implement AMICA's mixture of models, adaptive generalized-Gaussian source densities, or Newton updates, so they do not reproduce AMICA decompositions. pamica targets EEG/EMG analysts who want AMICA-quality decompositions inside a Python pipeline, users of GPU hardware who want faster runs than the CPU-only binary, and methodologists who need a transparent reference implementation to build on.

Implementation and validation

pamica provides a natural-gradient (Amari, 1998) expectation-maximization (EM) backend that ports the full AMICA algorithm: exact-EM mixture updates, a positive-definite Newton step (Palmer et al., 2008), symmetric zero-phase-component-analysis (ZCA) sphering, the five source-density families of the reference (generalized Gaussian, Gaussian, logistic, sub-Gaussian, and the extended-Infomax kurtosis switcher), a mixture of ICA models, and component sharing across models. It also computes mutual information reduction (MIR) and pairwise mutual information (PMI), the separation-quality metrics used to benchmark ICA algorithms (Delorme et al., 2012).

pamica's conformity with the reference binary is measured with two complementary metrics: Hungarian-matched component correlation and the Amari distance (Amari et al., 1996), a relabeling- and scale-invariant unmixing-matrix metric that needs no assignment step. Both implementations were run for AMICA's usual 2000 iterations with Newton off (`do_newton=0`) and otherwise-default parameters (settings transcribed between pamica's JSON and Fortran's native text format). Newton is disabled here to isolate the algorithm from initialization: with Newton on, independently seeded runs can settle a few under-determined components in different, equally likely optima, though a matched initialization recovers agreement (documentation). The single-model headline (Table 1) uses an external recording (OpenNeuro ds002718, 70 channels) with a data-adequacy ratio $k \approx 153$ (k = frames divided by squared channel count), well past the ~ 60 threshold where cross-backend agreement plateaus; the bundled 32-channel sample below ($k \approx 30$) gives a consistent Amari distance. Score functions and per-block sufficient statistics are exact to floating-point resolution against the literal Fortran expressions on the bundled sample. A mixture of ICA models is not partition-identifiable, so exact partition parity is the wrong bar for the multi-model case; it is instead assessed by whether the two implementations sample a similar distribution of solutions, across ensembles of 20 runs each (Figure 1). A permutation test finds no evidence that cross-implementation agreement is worse than Fortran's own run-to-run agreement. Multi-model log-likelihood distributions still differ slightly at a matched iteration budget (pamica needs about twice as many iterations to reach Fortran's mean), so full-likelihood similarity is not yet claimed.

Table 1: Single-model parity (external ds002718, $k \approx 153$; Amari and score-function checks on the bundled sample) and multi-model distributional similarity (bundled sample) of pamica with the Fortran reference. All values are means (sd = standard deviation, shown where relevant) over matched components (single-model) or within/cross-implementation pairs (190/400, multi-model).

Regime	Metric	Result (mean)
Single-model	Log-likelihood gap to Fortran (mean per-sample-channel log-likelihood, -3.6993 , $k \approx 153$)	within ~ 0.0005
Single-model	Conformity with Fortran	0.998 ($k \approx 153$) / 0.006 (bundled, $k \approx 30$)

Regime	Metric	Result (mean)
Single-model	Score functions (non-default families) and sufficient statistics	exact to floating-point resolution ($\sim 10^{-15}$)
Multi-model	Single-run magnitude (pamica-Fortran; Fortran-Fortran)	0.65; 0.64 (sd 0.05) / 0.163; 0.174 (sd 0.02)
Multi-model	Ensemble similarity, 20 runs each: mean difference, between — within-Fortran (permutation p)	+0.011 ($p = 0.96$) / -0.011 ($p > 0.999$)

Multi-model AMICA (n_models=2): pamica vs Fortran ensembles, real sample EEG

Dashed vertical lines mark each distribution's mean.




Figure 1: Multi-model solution-ensemble partition-correlation distributions (panel A) and log-likelihood distributions (panel B) for 20 pamica and 20 Fortran fits of the sample EEG; dashed lines mark each distribution's mean. The within-Fortran, within-pamica, and between-implementation correlation distributions overlap, so the single-run correlation reflects the estimator's intrinsic run-to-run spread rather than a gap to the reference. Panel B's apparent separation is a 0.009 log-likelihood gap on a ~ 0.035 axis.

79 All backends converge to the same single-model log-likelihood on real EEG (maximum pairwise
80 difference ~ 0.003). On Apple Silicon, MLX is the fastest backend and flat with channel count
81 (Table 2); PyTorch-MPS is never a win. Double-precision CUDA is the reproducible NVIDIA
82 path; native Fortran scales with CPU cores and, with enough cores pinned, can beat the
83 GPU on a larger, hotter host, though it does not match Apple's MLX on laptop hardware.
84 A data-size sweep (documentation) shows cross-backend component agreement rising with
85 frames per channel and plateauing near 0.98 once the decomposition is well-determined, where
86 two independent double-precision implementations agree at a mean of 0.995; single-precision
87 runs agree with float64 to four to five significant digits, so float64 stays the default for parity.

Table 2: Single-model throughput on real 70-channel EEG ($n_{\text{mix}}=3$, $\text{pdftype}=0$, $\text{block_size}=512$; warm, minimum of repeated runs). CPU, MPS, and MLX on Apple Silicon; CUDA on a separate NVIDIA RTX 4090 (float32 comparable, ~ 36 ms). The two native-Fortran rows are from a separate core-count sweep (documentation) at each backend's plateau; the other CPU rows use platform-default threads and are not core-matched to Fortran. Unlike the correctness comparison, this benchmark uses external data (OpenNeuro ds002718, one subject so far) and specific GPU hardware.

Backend (device)	Precision	ms / iteration
MLX (Apple GPU)	float32	25
CUDA (NVIDIA RTX 4090)	float64	39
Native Fortran (Intel Core i9-13900K, 24 cores)	float64	30
PyTorch CPU (Apple Silicon)	float64	193
Native Fortran (Apple Silicon, 8 cores)	float64	70
PyTorch MPS (Apple GPU)	float32	255
NumPy (reference, Apple Silicon)	float64	622

88 The correctness harness never uses synthetic data; the multi-model and score-function checks
 89 need no external download (bundled sample only). The full performance tables, per-run
 90 Amari-distance detail, data-size sweep, and step-by-step reproduction commands are in the
 91 documentation (<https://eeglab.org/pAMICA/guides/validation/>).

92 State of the field

93 *pamica* complements rather than replaces EEGLAB (Delorme & Makeig, 2004) and its Fortran
 94 AMICA plugin: it shares EEGLAB's output format, adds a Python API with GPU support,
 95 and runs the reference Fortran itself through a bundled dependency-free native build (no Intel
 96 Math Kernel Library or MPI runtime). Two other Python AMICA reimplementations have
 97 appeared (Esmaeili, 2025; Herforth, 2026), both of which provide MNE-Python-compatible
 98 objects; *pamica* offers a scikit-learn-style array API, byte-identical EEGLAB I/O, an MLX
 99 backend for Apple GPUs, and an optional MNE-Python wrapper. What sets *pamica* apart is
 100 the depth of its Fortran-parity validation (Table 1): source-density score functions bit-exact to
 101 floating-point resolution against the literal Fortran expressions, and a distributional-similarity
 102 framework for the non-identifiable multi-model case.

103 Acknowledgements

104 We thank Jason Palmer and Ken Kreutz-Delgado, co-developers of AMICA, for the reference
 105 implementation. We also thank the EEGLAB community for the tools and sample data used to
 106 validate this work. Two of the authors are original developers of the methods *pamica* builds on:
 107 S.M. co-developed the AMICA algorithm (Palmer et al., 2012) and A.D. is a lead developer of
 108 EEGLAB (Delorme & Makeig, 2004). This work was supported by The Swartz Foundation (Old
 109 Field, NY) to the Swartz Center for Computational Neuroscience and by National Institutes of
 110 Health grant R01-NS047293 (to A.D. and S.M.).

111 References

- 112 Ablin, P., Cardoso, J.-F., & Gramfort, A. (2018). Faster independent component analysis
 113 by preconditioning with hessian approximations. *IEEE Transactions on Signal Processing*,
 114 66(15), 4040–4053. <https://doi.org/10.1109/TSP.2018.2844203>
- 115 Amari, S.-I. (1998). Natural gradient works efficiently in learning. *Neural Computation*, 10(2),
 116 251–276. <https://doi.org/10.1162/089976698300017746>

- Amari, S.-I., Cichocki, A., & Yang, H. H. (1996). A new learning algorithm for blind signal separation. *Advances in Neural Information Processing Systems*, 8, 757–763.
- Bell, A. J., & Sejnowski, T. J. (1995). An information-maximization approach to blind separation and blind deconvolution. *Neural Computation*, 7(6), 1129–1159. <https://doi.org/10.1162/neco.1995.7.6.1129>
- Delorme, A., & Makeig, S. (2004). EEGLAB: An open source toolbox for analysis of single-trial EEG dynamics including independent component analysis. *Journal of Neuroscience Methods*, 134(1), 9–21. <https://doi.org/10.1016/j.jneumeth.2003.10.009>
- Delorme, A., Palmer, J., Onton, J., Oostenveld, R., & Makeig, S. (2012). Independent EEG sources are dipolar. *PLoS ONE*, 7(2), e30135. <https://doi.org/10.1371/journal.pone.0030135>
- Esmaili, S. (2025). *Amica: Native python implementation of AMICA for MNE-Python and scientific EEG workflows*. <https://github.com/snesmaeli/amica>.
- Gramfort, A., Luessi, M., Larson, E., Engemann, D. A., Strohmeier, D., Brodbeck, C., Goj, R., Jas, M., Brooks, T., Parkkonen, L., & Hämäläinen, M. (2013). MEG and EEG data analysis with MNE-Python. *Frontiers in Neuroscience*, 7, 267. <https://doi.org/10.3389/fnins.2013.00267>
- Hannun, A., Digani, J., Katharopoulos, A., & Collobert, R. (2023). *MLX: Efficient and flexible machine learning on apple silicon*. <https://github.com/ml-explore/mlx>.
- Harris, C. R., Millman, K. J., Walt, S. J. van der, Gommers, R., Virtanen, P., Cournapeau, D., Wieser, E., Taylor, J., Berg, S., Smith, N. J., & others. (2020). Array programming with NumPy. *Nature*, 585(7825), 357–362. <https://doi.org/10.1038/s41586-020-2649-2>
- Herforth, J. (2026). *Pyamica: PyTorch AMICA, adaptive mixture independent component analysis*. <https://github.com/DerAndereJohannes/pyamica>.
- Hyvärinen, A., & Oja, E. (2000). Independent component analysis: Algorithms and applications. *Neural Networks*, 13(4-5), 411–430. [https://doi.org/10.1016/S0893-6080\(00\)00026-5](https://doi.org/10.1016/S0893-6080(00)00026-5)
- Lee, T.-W., Girolami, M., & Sejnowski, T. J. (1999). Independent component analysis using an extended infomax algorithm for mixed subgaussian and supergaussian sources. *Neural Computation*, 11(2), 417–441. <https://doi.org/10.1162/089976699300016719>
- Palmer, J. A., Kreutz-Delgado, K., & Makeig, S. (2012). *AMICA: An adaptive mixture of independent component analyzers with shared components*. Swartz Center for Computational Neuroscience, University of California San Diego. https://scn.ucsd.edu/~jason/amica_a.pdf
- Palmer, J. A., Kreutz-Delgado, K., Rao, B. D., & Makeig, S. (2008). Newton method for the ICA mixture model. *2008 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 1805–1808. <https://doi.org/10.1109/ICASSP.2008.4517982>
- Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z., Gimelshein, N., Antiga, L., & others. (2019). PyTorch: An imperative style, high-performance deep learning library. *Advances in Neural Information Processing Systems*, 32, 8024–8035.
- Pion-Tonachini, L., Kreutz-Delgado, K., & Makeig, S. (2019). ICLabel: An automated electroencephalographic independent component classifier, dataset, and website. *NeuroImage*, 198, 181–197. <https://doi.org/10.1016/j.neuroimage.2019.05.026>
- Virtanen, P., Gommers, R., Oliphant, T. E., Haberland, M., Reddy, T., Cournapeau, D., Burovski, E., Peterson, P., Weckesser, W., Bright, J., & others. (2020). SciPy 1.0: Fundamental algorithms for scientific computing in python. *Nature Methods*, 17(3), 261–272. <https://doi.org/10.1038/s41592-019-0686-2>