

pyAMICA: GPU-accelerated Adaptive Mixture Independent Component Analysis in Python with Fortran parity

Seyed Yahya Shirazi¹, Arnaud Delorme^{1,2}, and Scott Makeig¹

¹ Swartz Center for Computational Neuroscience, Institute for Neural Computation, University of California San Diego, USA ² Centre de Recherche Cerveau et Cognition (CerCo), CNRS, University of Toulouse, France ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Independent Component Analysis (ICA) is a standard method for separating electroencephalography (EEG) and electromyography (EMG) recordings into maximally independent sources, which isolates brain, muscle, and artifact activity for downstream analysis. Adaptive Mixture ICA (AMICA) generalizes single-model ICA to a mixture of ICA models with adaptive source densities, and produces the most dipolar (and thus most physiologically interpretable) component decompositions of EEG among the widely used algorithms benchmarked by Delorme et al. (2012). Its reference implementation, however, is a Fortran program parallelized with the Message Passing Interface (MPI) and distributed as a compiled binary driven from MATLAB/EEGLAB, which is difficult to install, runs only on the CPU, and is not usable from a Python scientific workflow.

pyAMICA is a Python implementation of AMICA that reproduces the reference Fortran results within numerical tolerance while running on the CPU, NVIDIA GPUs (CUDA), and Apple GPUs (Apple's MLX array framework (Hannun et al., 2023)). It is a complete NumPy/PyTorch reimplementation, not a wrapper around the Fortran binary, built on PyTorch (Paszke et al., 2019), NumPy (Harris et al., 2020), and SciPy (Virtanen et al., 2020) (Python \$ \$3.12, PyTorch \$ \$2.12; tested on Apple Silicon/ARM64 and x86_64/CUDA Linux), and exposes a scikit-learn-style estimator under a BSD-3-Clause license. In double precision it reproduces the reference score-function algebra to machine precision; it also runs in single precision (float32), which the CPU-only binary does not offer and which is numerically faithful (agreeing with double precision to four to five significant digits) yet required to use Apple GPUs, which have no float64 and host the fastest backend (MLX). pyAMICA writes output in the binary format that EEGLAB's AMICA loader reads: a single-model output file is byte-identical in layout to a native AMICA file and needs no manual re-interpretation, and multi-model output round-trips through the same loader. Correctness is defined as parity with the Fortran reference for the single-model case and, because multi-model AMICA is not partition-identifiable, as distributional equivalence for the multi-model case; both are validated on real EEG against the reference binary. A minimal example is a three-line scikit-learn-style call, `AMICA(n_models=1, n_mix=3).fit(X)` on a (channels, samples) array (full workflow in the README). The software is at <https://github.com/scen/pyAMICA> (archived at doi:10.5281/zenodo.21312148).

Statement of need

AMICA yields components that are well suited to equivalent-dipole source localization and to automated classification of independent components (Pion-Tonachini et al., 2019), and it separates EEG more effectively than most alternatives (Delorme et al., 2012; Palmer et

42 [al., 2012](#)). The reference implementation is the original Fortran code: it depends on an MPI
43 toolchain and precompiled binaries that are awkward to build across platforms, it cannot use a
44 GPU, and it is invoked as an external process from MATLAB rather than called from Python.
45 As neuroimaging analysis has moved toward Python, for example MNE-Python ([Gramfort et](#)
46 [al., 2013](#)), an AMICA that runs natively in Python and on a GPU, and that is validated to
47 reproduce the Fortran reference numerically, is needed for modern pipelines and for studying
48 the algorithm in an open codebase.

49 General-purpose Python ICA implementations do not fill this gap. scikit-learn and MNE-
50 Python provide FastICA ([Hyvärinen & Oja, 2000](#)) and Infomax ([Bell & Sejnowski, 1995](#); [Lee](#)
51 [et al., 1999](#)), and Picard ([Ablin et al., 2018](#)) provides a faster maximum-likelihood ICA, but
52 none implement AMICA's mixture of models, adaptive generalized-Gaussian source densities,
53 or Newton updates, and so they do not reproduce AMICA decompositions. pyAMICA targets
54 researchers who need AMICA specifically: EEG/EMG analysts who want AMICA-quality
55 decompositions inside a Python pipeline, users of GPU hardware who want faster runs than
56 the CPU-only binary, and methodologists who need a transparent reference implementation to
57 inspect and build on.

58 Implementation and validation

59 pyAMICA provides a natural-gradient ([Amari, 1998](#)) expectation-maximization backend that
60 ports the full AMICA algorithm: exact-EM mixture updates, a positive-definite Newton step
61 ([Palmer et al., 2008](#)), symmetric zero-phase-component-analysis (ZCA) sphering, the five
62 source-density families of the reference (generalized Gaussian, Gaussian, logistic, sub-Gaussian,
63 and the extended-Infomax kurtosis switcher), a mixture of ICA models, and component sharing
64 across models.

65 pyAMICA's conformity with the reference binary is measured on one exemplar real EEG recording
66 (the bundled EEGLAB tutorial dataset: 32 channels, 30504 samples at 128 Hz, ~238 s;
67 Table 1) – a multi-subject, multi-dataset validation is planned future work – running both
68 implementations for AMICA's usual 2000 iterations with otherwise-default parameters, with two
69 complementary metrics: Hungarian-matched component correlation, and the Amari distance
70 ([Amari et al., 1996](#)), a standard unmixing-matrix comparison metric that needs no assignment
71 step since it is permutation- and scale-invariant by construction. Both agree closely for the
72 single model (mean values; Table 1), and the source-density score functions and per-block
73 sufficient statistics are exact to floating-point resolution against the literal Fortran expressions.
74 A mixture of ICA models is not partition-identifiable, so exact partition parity is the wrong bar
75 for the multi-model case; it is instead assessed by distributional equivalence across ensembles
76 of 20 runs each ([Figure 1](#)). Both metrics again agree: a run-level permutation test, which
77 permutes the 40 runs as intact units to respect the dependence among pairwise values, finds
78 no evidence that cross-implementation agreement is worse than Fortran's own run-to-run
79 agreement. The single-run values are therefore intrinsic estimator spread rather than a shortfall.
80 Equivalence is claimed for the partition structure; the multi-model log-likelihood distributions
81 still differ slightly at a matched 100-iteration budget (pyAMICA reaches Fortran's mean with
82 about twice as many iterations), so full-likelihood equivalence is not yet claimed. Per-run
83 detail for both metrics is in the documentation.

Table 1: Single-model parity and multi-model distributional equivalence of pyAMICA with the Fortran reference on the bundled sample EEG. All values are means (sd shown where relevant) over matched components (single-model) or over the 190 pairwise ensemble comparisons (multi-model).

Regime	Metric	Result (mean; correlation / Amari distance)
Single-model	Log-likelihood gap to Fortran (mean per-sample-channel LL, -3.4018)	within ~ 0.005
Single-model	Conformity with Fortran	0.997 / 0.006
Single-model	Score functions and sufficient statistics	exact to floating-point resolution ($\sim 10^{-15}$)
Multi-model	Single-run magnitude (pyAMICA-Fortran; Fortran-Fortran)	0.65; 0.64 (sd 0.05) / 0.163; 0.174 (sd 0.02)
Multi-model	Ensemble equivalence, 20 runs each: mean difference, between — within-Fortran (permutation p)	$+0.011$ ($p = 0.96$) / -0.011 ($p > 0.999$)

Multi-model AMICA ($n_{\text{models}}=2$): pyAMICA vs Fortran ensembles, real sample EEG

Dashed vertical lines mark each distribution's mean.




Figure 1: Multi-model solution-ensemble partition-correlation distributions (panel A) and log-likelihood distributions (panel B) for 20 pyAMICA and 20 Fortran fits of the sample EEG; dashed lines mark each distribution's mean. The within-Fortran, within-pyAMICA, and between-implementation correlation distributions overlap, so the single-run correlation reflects the estimator's intrinsic run-to-run spread rather than a gap to the reference. Panel B's apparent separation is a mean log-likelihood gap of only 0.009 on an axis spanning ~ 0.02 .

84 All backends converge to the same single-model log-likelihood on real EEG (maximum pairwise
85 difference ~ 0.003). On Apple Silicon the MLX backend is the fastest option and is roughly flat
86 with channel count (15-25 ms per iteration from 16 to 70 channels; see the documentation),
87 about eight times faster than double-precision multithreaded CPU; PyTorch-MPS is not a win
88 (at or worse than the CPU). On NVIDIA hardware double-precision CUDA is the reproducible
89 path and is overhead-bound at EEG scale, so single precision gives it little additional speedup
90 (Table 2). Native Fortran itself scales with CPU cores, unlike the CPU backends above:
91 with enough cores pinned it is competitive with, or faster than, the GPU it is compared
92 against on each machine (Table 2), though only by dedicating most cores of a much larger,
93 hotter host than a laptop GPU. A data-size sweep further shows cross-backend component
94 equivalence rising with frames per channel and plateauing near 0.98 once the decomposition is
95 well-determined, where two independent double-precision implementations (native Fortran and
96 PyTorch-CUDA) agree at a mean of 0.995; single-precision runs agree with double precision to
97 seven significant digits, so double precision remains the default for parity.

Table 2: Single-model throughput on real 70-channel EEG ($n_{\text{mix}}=3$, $\text{pdftype}=0$, $\text{block_size}=512$; warm, minimum of repeated runs). CPU, MPS, and MLX on Apple Silicon; CUDA on a separate NVIDIA RTX 4090; CUDA float32 is comparable (~ 36 ms). The two native-Fortran rows are from a separate core-count sweep (documentation) on the same two machines, at the core count where each backend's throughput levels off; the other CPU rows above use the platform default thread count, so they are not core-matched to Fortran. Unlike the correctness comparison, this benchmark uses external data (OpenNeuro ds002718, one subject so far) and specific GPU hardware.

Backend (device)	Precision	ms / iteration
MLX (Apple GPU)	float32	25
CUDA (NVIDIA RTX 4090)	float64	39
Native Fortran (Intel Core i9-13900K, 24 cores)	float64	30
PyTorch CPU (Apple Silicon)	float64	193
Native Fortran (Apple Silicon, 8 cores)	float64	70
PyTorch MPS (Apple GPU)	float32	255
NumPy (reference, Apple Silicon)	float64	622

98 The correctness harness compares pyAMICA against Fortran with two metrics, Hungarian-
99 matched component correlation and Amari distance, and uses only the bundled real sample
100 EEG and Fortran binary, with no external download and no synthetic data. The full per-channel
101 and multi-model performance tables, the per-run Amari-distance detail, and the data-size
102 sweep, along with the step-by-step commands to reproduce every number here, are in the
103 documentation (<https://eeglab.org/pyAMICA/guides/validation/>).

104 State of the field

105 pyAMICA complements rather than replaces EEGLAB (Delorme & Makeig, 2004) and its Fortran
106 AMICA plugin: it reads and writes the same output format, so results move between the two,
107 while adding GPU support and a Python API. Two other Python AMICA reimplementations
108 have appeared (Esmaeili, 2025; Herforth, 2026), both of which provide MNE-Python-compatible
109 objects; pyAMICA instead offers a scikit-learn-style array API and byte-identical EEGLAB I/O,
110 and does not yet ship an MNE-Python wrapper. What distinguishes pyAMICA is the rigor and
111 scope of its Fortran-parity validation, beyond what either alternative publishes: source-density
112 score functions exact to floating-point resolution ($\sim 10^{-15}$), single-model component correlation
113 and Amari distance against the reference, and a distributional-equivalence framework for the
114 non-identifiable multi-model case, together with byte-identical EEGLAB output and an MLX
115 backend for Apple GPUs.

Acknowledgements

We thank Jason Palmer and Ken Kreutz-Delgado, co-developers of AMICA, for the reference implementation, and the EEGLAB community for the tools and sample data used to validate this work. Two of the authors are original developers of the methods pyAMICA builds on: S.M. co-developed the AMICA algorithm (Palmer et al., cited above) and A.D. is a lead developer of EEGLAB (Delorme & Makeig, 2004). This work was supported by The Swartz Foundation (Old Field, NY) to the Swartz Center for Computational Neuroscience and by National Institutes of Health grant R01-NS047293 (to A.D. and S.M.).

References

- Ablin, P., Cardoso, J.-F., & Gramfort, A. (2018). Faster independent component analysis by preconditioning with hessian approximations. *IEEE Transactions on Signal Processing*, 66(15), 4040–4053. <https://doi.org/10.1109/TSP.2018.2844203>
- Amari, S.-I. (1998). Natural gradient works efficiently in learning. *Neural Computation*, 10(2), 251–276. <https://doi.org/10.1162/089976698300017746>
- Amari, S.-I., Cichocki, A., & Yang, H. H. (1996). A new learning algorithm for blind signal separation. *Advances in Neural Information Processing Systems*, 8, 757–763.
- Bell, A. J., & Sejnowski, T. J. (1995). An information-maximization approach to blind separation and blind deconvolution. *Neural Computation*, 7(6), 1129–1159. <https://doi.org/10.1162/neco.1995.7.6.1129>
- Delorme, A., & Makeig, S. (2004). EEGLAB: An open source toolbox for analysis of single-trial EEG dynamics including independent component analysis. *Journal of Neuroscience Methods*, 134(1), 9–21. <https://doi.org/10.1016/j.jneumeth.2003.10.009>
- Delorme, A., Palmer, J., Onton, J., Oostenveld, R., & Makeig, S. (2012). Independent EEG sources are dipolar. *PLoS ONE*, 7(2), e30135. <https://doi.org/10.1371/journal.pone.0030135>
- Esmaeili, S. (2025). *Amica: Native python implementation of AMICA for MNE-Python and scientific EEG workflows*. <https://github.com/snesmaeili/amica>.
- Gramfort, A., Luessi, M., Larson, E., Engemann, D. A., Strohmeier, D., Brodbeck, C., Goj, R., Jas, M., Brooks, T., Parkkonen, L., & Hämäläinen, M. (2013). MEG and EEG data analysis with MNE-Python. *Frontiers in Neuroscience*, 7, 267. <https://doi.org/10.3389/fnins.2013.00267>
- Hannun, A., Digani, J., Katharopoulos, A., & Collobert, R. (2023). *MLX: Efficient and flexible machine learning on apple silicon*. <https://github.com/ml-explore/mlx>.
- Harris, C. R., Millman, K. J., Walt, S. J. van der, Gommers, R., Virtanen, P., Cournapeau, D., Wieser, E., Taylor, J., Berg, S., Smith, N. J., & others. (2020). Array programming with NumPy. *Nature*, 585(7825), 357–362. <https://doi.org/10.1038/s41586-020-2649-2>
- Herforth, J. (2026). *Pyamica: PyTorch AMICA, adaptive mixture independent component analysis*. <https://github.com/DerAndereJohannes/pyamica>.
- Hyvärinen, A., & Oja, E. (2000). Independent component analysis: Algorithms and applications. *Neural Networks*, 13(4-5), 411–430. [https://doi.org/10.1016/S0893-6080\(00\)00026-5](https://doi.org/10.1016/S0893-6080(00)00026-5)
- Lee, T.-W., Girolami, M., & Sejnowski, T. J. (1999). Independent component analysis using an extended infomax algorithm for mixed subgaussian and supergaussian sources. *Neural Computation*, 11(2), 417–441. <https://doi.org/10.1162/089976699300016719>
- Palmer, J. A., Kreutz-Delgado, K., & Makeig, S. (2012). *AMICA: An adaptive*

- 160 *mixture of independent component analyzers with shared components.* Swartz
161 Center for Computational Neuroscience, University of California San Diego.
162 https://sccn.ucsd.edu/~jason/amica_a.pdf
- 163 Palmer, J. A., Kreutz-Delgado, K., Rao, B. D., & Makeig, S. (2008). Newton method for the
164 ICA mixture model. *2008 IEEE International Conference on Acoustics, Speech and Signal*
165 *Processing (ICASSP)*, 1805–1808. <https://doi.org/10.1109/ICASSP.2008.4517982>
- 166 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin,
167 Z., Gimelshein, N., Antiga, L., & others. (2019). PyTorch: An imperative style, high-
168 performance deep learning library. *Advances in Neural Information Processing Systems*, 32,
169 8024–8035.
- 170 Pion-Tonachini, L., Kreutz-Delgado, K., & Makeig, S. (2019). ICLabel: An automated
171 electroencephalographic independent component classifier, dataset, and website.
172 *NeuroImage*, 198, 181–197. <https://doi.org/10.1016/j.neuroimage.2019.05.026>
- 173 Virtanen, P., Gommers, R., Oliphant, T. E., Haberland, M., Reddy, T., Cournapeau, D.,
174 Burovski, E., Peterson, P., Weckesser, W., Bright, J., & others. (2020). SciPy 1.0:
175 Fundamental algorithms for scientific computing in python. *Nature Methods*, 17(3),
176 261–272. <https://doi.org/10.1038/s41592-019-0686-2>

DRAFT